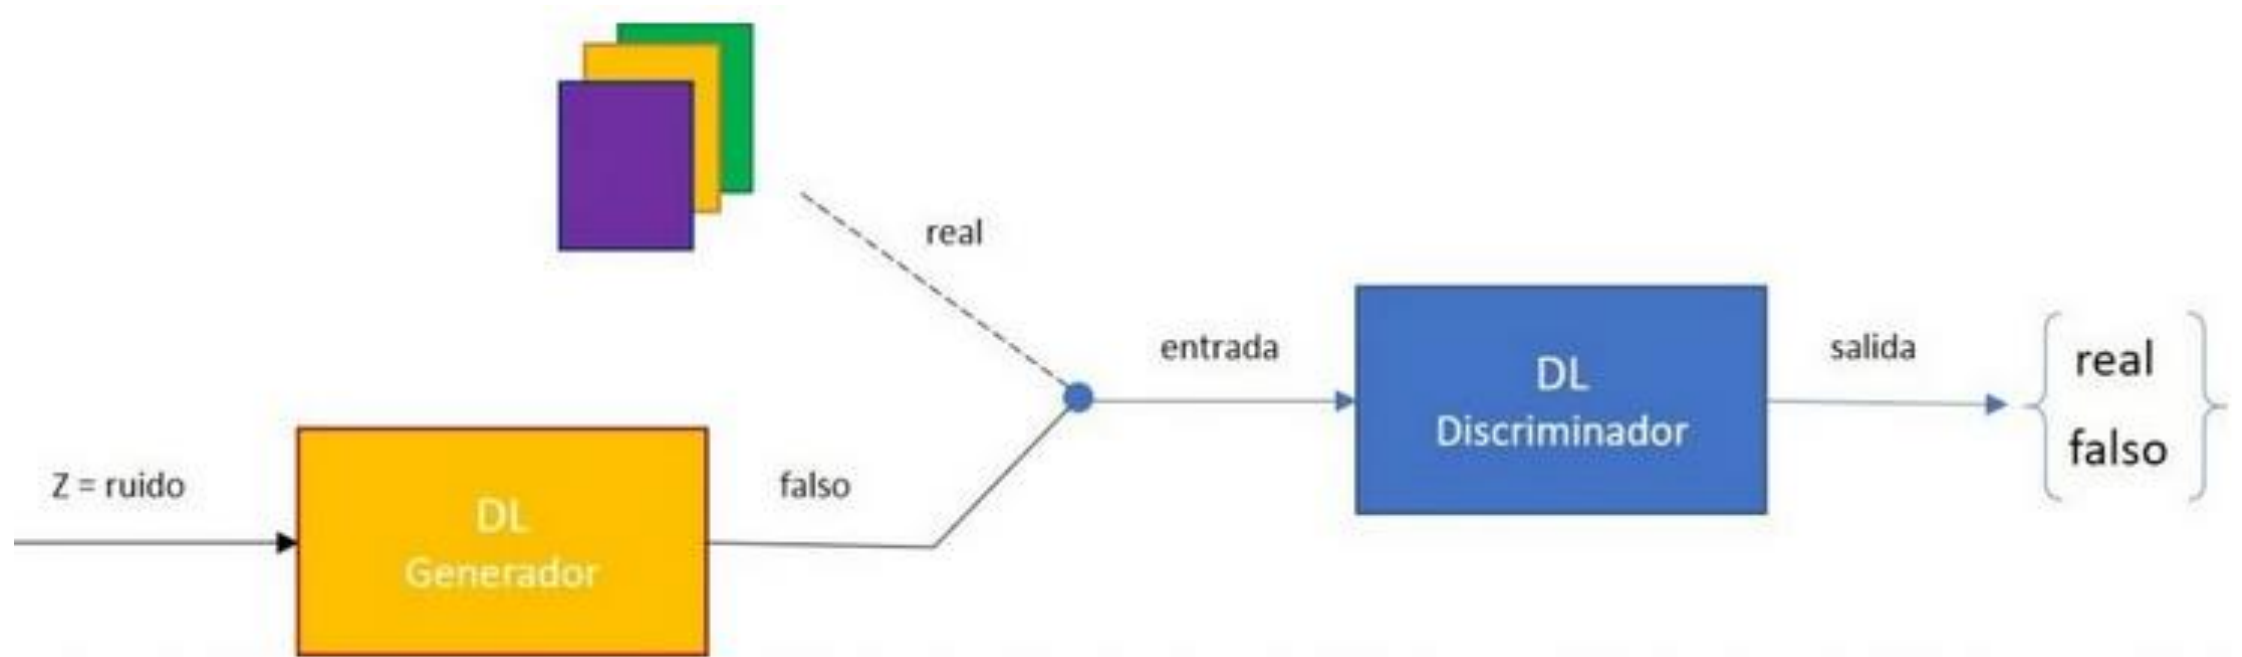
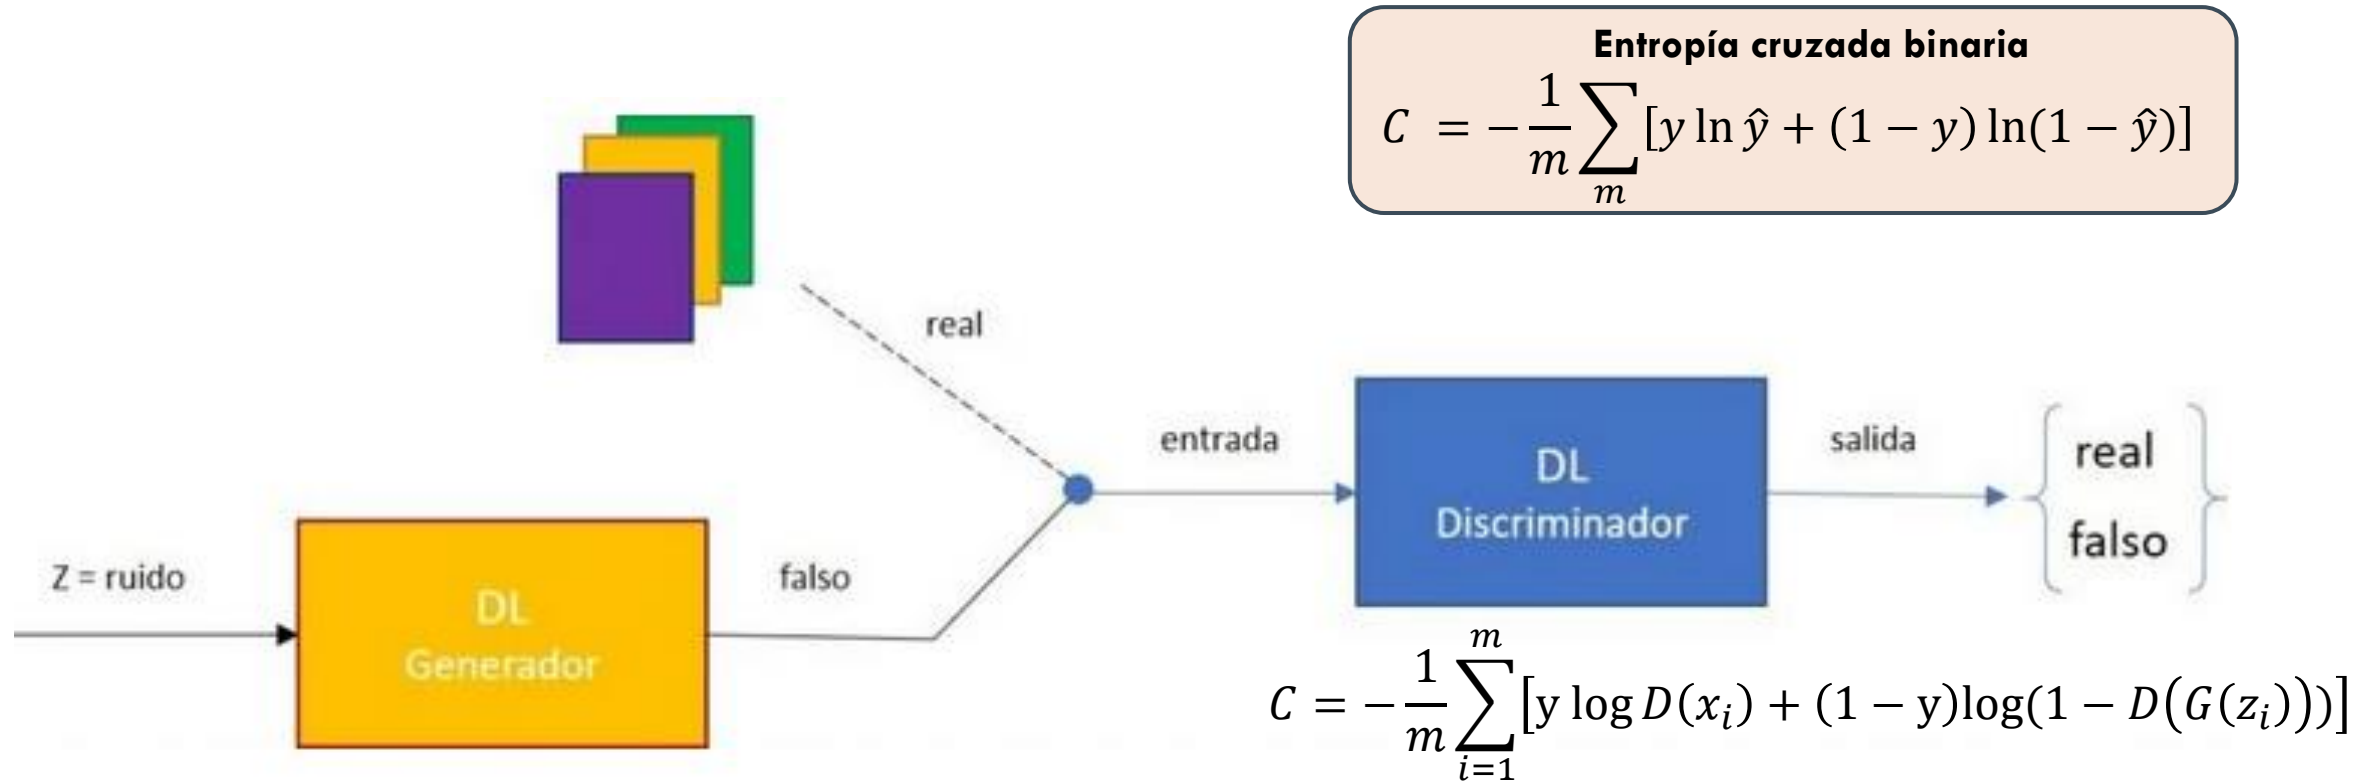


# Redes generativas adversarias (GAN)



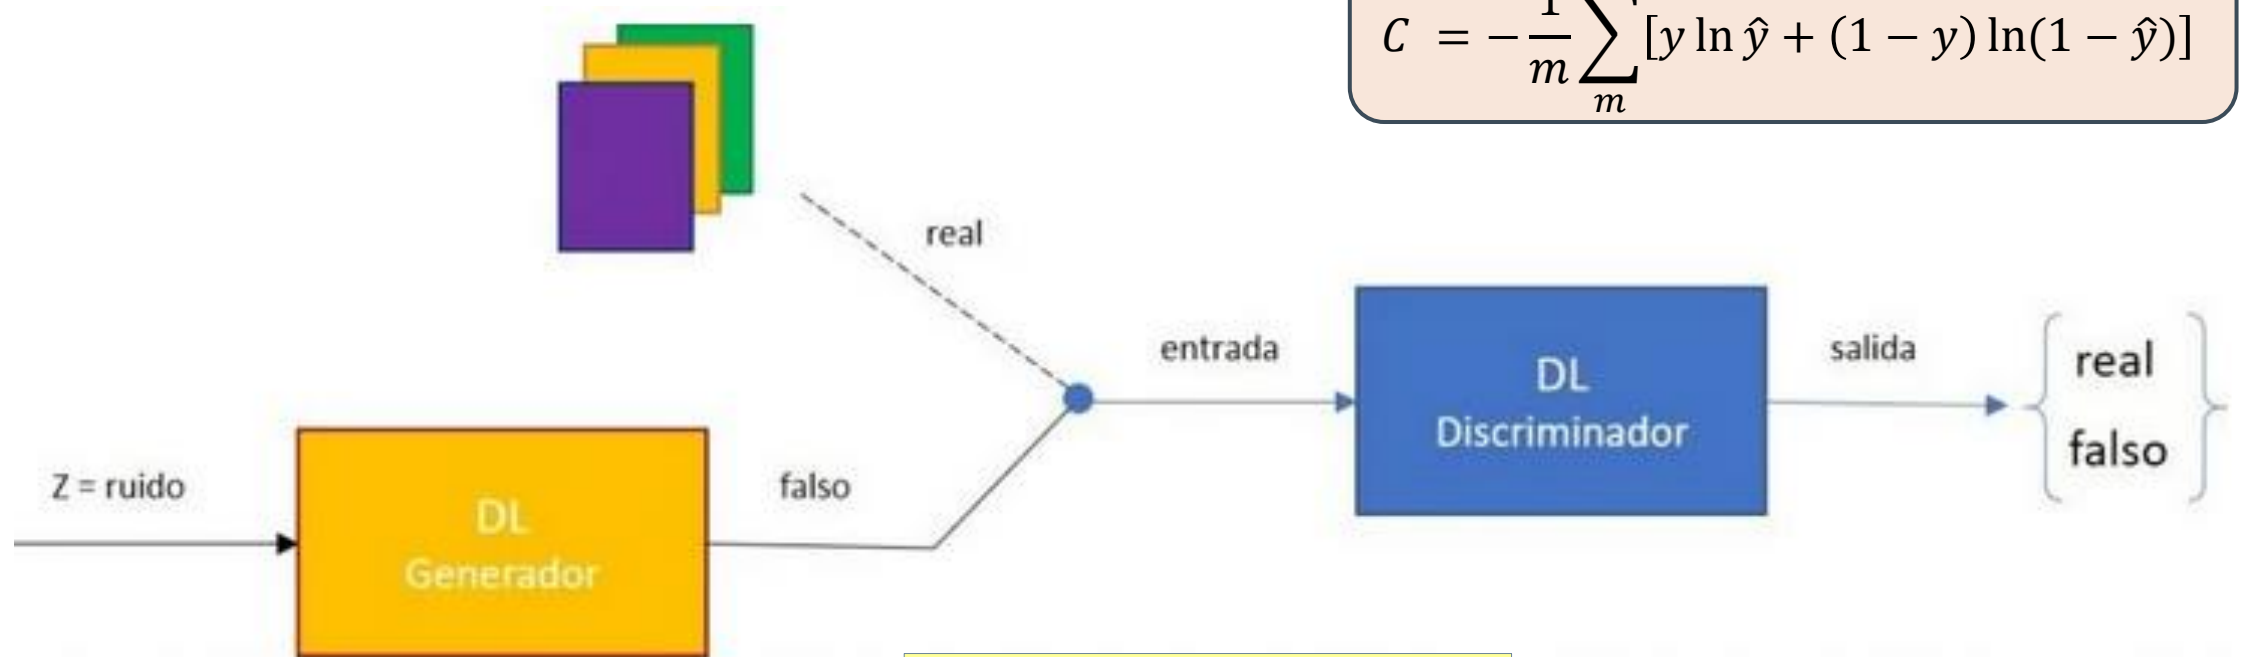
# Redes generativas adversarias (GAN)



# Redes generativas adversarias (GAN)

**Entropía cruzada binaria**

$$C = -\frac{1}{m} \sum_m [y \ln \hat{y} + (1 - y) \ln(1 - \hat{y})]$$



$$C = -\frac{1}{m} \sum_{i=1}^m [\log(D(G(z_i)))]$$

Considera sólo cuando el discriminador clasifica incorrectamente las muestras generadas

# Discriminador

```
def build_discriminator(img_shape):  
    model = Sequential()  
    model.add(Flatten(input_shape=img_shape))  
    model.add(Dense(512))  
    model.add(LeakyReLU(alpha=0.2))  
    model.add(Dense(256))  
    model.add(LeakyReLU(alpha=0.2))  
    model.add(Dense(1, activation='sigmoid'))  
    return model
```

```
img_shape = (28, 28, 1)
```

*# Crear y compilar el discriminador*

```
discriminator = build_discriminator(img_shape)  
discriminator.compile(loss='binary_crossentropy', optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
```

Layer (type)	Output Shape	Param #
=====		
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 512)	401920
leaky_re_lu (LeakyReLU)	(None, 512)	0
dense_1 (Dense)	(None, 256)	131328
leaky_re_lu_1 (LeakyReLU)	(None, 256)	0
dense_2 (Dense)	(None, 1)	257
=====		
Total params: 533,505		
Trainable params: 533,505		
Non-trainable params: 0		

# Generator

```
def build_generator(latent_dim):  
    model = Sequential()  
    model.add(Dense(256, input_dim=latent_dim))  
    model.add(LeakyReLU(alpha=0.2))  
    model.add(BatchNormalization(momentum=0.8))  
    model.add(Dense(512))  
    model.add(LeakyReLU(alpha=0.2))  
    model.add(BatchNormalization(momentum=0.8))  
    model.add(Dense(1024))  
    model.add(LeakyReLU(alpha=0.2))  
    model.add(BatchNormalization(momentum=0.8))  
    model.add(Dense(28 * 28 * 1, activation='tanh'))  
    model.add(Reshape((28, 28, 1)))  
    return model
```

```
latent_dim = 100
```

```
# Crear el generador
```

```
generator = build_generator(latent_dim)
```

Layer (type)	Output Shape	Param #
=====		
dense_3 (Dense)	(None, 256)	25856
leaky_re_lu_2 (LeakyReLU)	(None, 256)	0
batch_normalization	(None, 256)	1024
dense_4 (Dense)	(None, 512)	131584
leaky_re_lu_3 (LeakyReLU)	(None, 512)	0
batch_normalization_1	(None, 512)	2048
dense_5 (Dense)	(None, 1024)	525312
leaky_re_lu_4 (LeakyReLU)	(None, 1024)	0
batch_normalization_2	(None, 1024)	4096
dense_6 (Dense)	(None, 784)	803600
reshape (Reshape)	(None, 28, 28, 1)	0
=====		
Total params: 1,493,520		
Trainable params: 1,489,936		
Non-trainable params: 3,584		

# Modelo combinado. Generador+Discriminador

```
latent_dim = 100
```

```
# El generador toma ruido como entrada y genera imágenes
```

```
z = Input(shape=(latent_dim,))
```

```
img = generator(z)
```

```
# En el modelo combinado solo se entrena el generador
```

```
discriminator.trainable = False
```

```
# Rta del discriminador
```

```
valid = discriminator(img)
```

```
# Modelo combinado (stacked generador y discriminador)
```

```
combined = Model(z, valid)
```

```
combined.compile(loss='binary_crossentropy', optimizer=Adam(0.0002, 0.5))
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 100)]	0
sequential_1 (Sequential)	(None, 28, 28, 1)	1493520
sequential (Sequential)	(None, 1)	533505

=====  
Total params: 2,027,025  
Trainable params: 1,489,936  
Non-trainable params: 537,089

# Entrenamiento del modelo

for epoch in range(epochs):

*Discriminador*

```
idx = np.random.randint(0, X_train.shape[0], batch_size)
imgs = X_train[idx]

noise = np.random.normal(0, 1, (batch_size, latent_dim))
gen_imgs = generator.predict(noise)

d_loss_real = discriminator.train_on_batch(imgs, valid)
d_loss_fake = discriminator.train_on_batch(gen_imgs, fake)
d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
```

*Seleccionar un conjunto aleatorio de imágenes reales*

```
noise = np.random.normal(0, 1, (batch_size, latent_dim))
g_loss = combined.train_on_batch(noise, valid)
```

# Entrenamiento del modelo

for epoch in range(epochs):

*Discriminador*

```
idx = np.random.randint(0, X_train.shape[0], batch_size)
imgs = X_train[idx]
```

```
noise = np.random.normal(0, 1, (batch_size, latent_dim))
gen_imgs = generator.predict(noise)
```

```
d_loss_real = discriminator.train_on_batch(imgs, valid)
d_loss_fake = discriminator.train_on_batch(gen_imgs, fake)
d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
```

*Generar un conjunto de imágenes falsas*

```
noise = np.random.normal(0, 1, (batch_size, latent_dim))
g_loss = combined.train_on_batch(noise, valid)
```



# Entrenamiento del modelo

for epoch in range(epochs):

*Discriminador*

```
idx = np.random.randint(0, X_train.shape[0], batch_size)
imgs = X_train[idx]

noise = np.random.normal(0, 1, (batch_size, latent_dim))
gen_imgs = generator.predict(noise)

d_loss_real = discriminator.train_on_batch(imgs, valid)
d_loss_fake = discriminator.train_on_batch(gen_imgs, fake)
d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
```

*Entrenar el discriminador*

```
noise = np.random.normal(0, 1, (batch_size, latent_dim))
g_loss = combined.train_on_batch(noise, valid)
```

# Entrenamiento del modelo

for epoch in range(epochs):

*Discriminador*

```
idx = np.random.randint(0, X_train.shape[0], batch_size)
imgs = X_train[idx]

noise = np.random.normal(0, 1, (batch_size, latent_dim))
gen_imgs = generator.predict(noise)

d_loss_real = discriminator.train_on_batch(imgs, valid)
d_loss_fake = discriminator.train_on_batch(gen_imgs, fake)
d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
```


```
noise = np.random.normal(0, 1, (batch_size, latent_dim))
g_loss = combined.train_on_batch(noise, valid)
```

*Entrena el generador para engañar al discriminador*

**GAN\_MNIST.ipynb**

# Evolución del entrenamiento


## □ Época 0



*GAN\_MNIST.ipynb*

# Evolución del entrenamiento


## □ Época 400



*GAN\_MNIST.ipynb*

# Evolución del entrenamiento


## □ Época 600



*GAN\_MNIST.ipynb*

# Evolución del entrenamiento


## □ Época 800



*GAN\_MNIST.ipynb*

# Evolución del entrenamiento

## □ Época 1600



*GAN\_MNIST.ipynb*

# Evolución del entrenamiento


## □ Época 3200





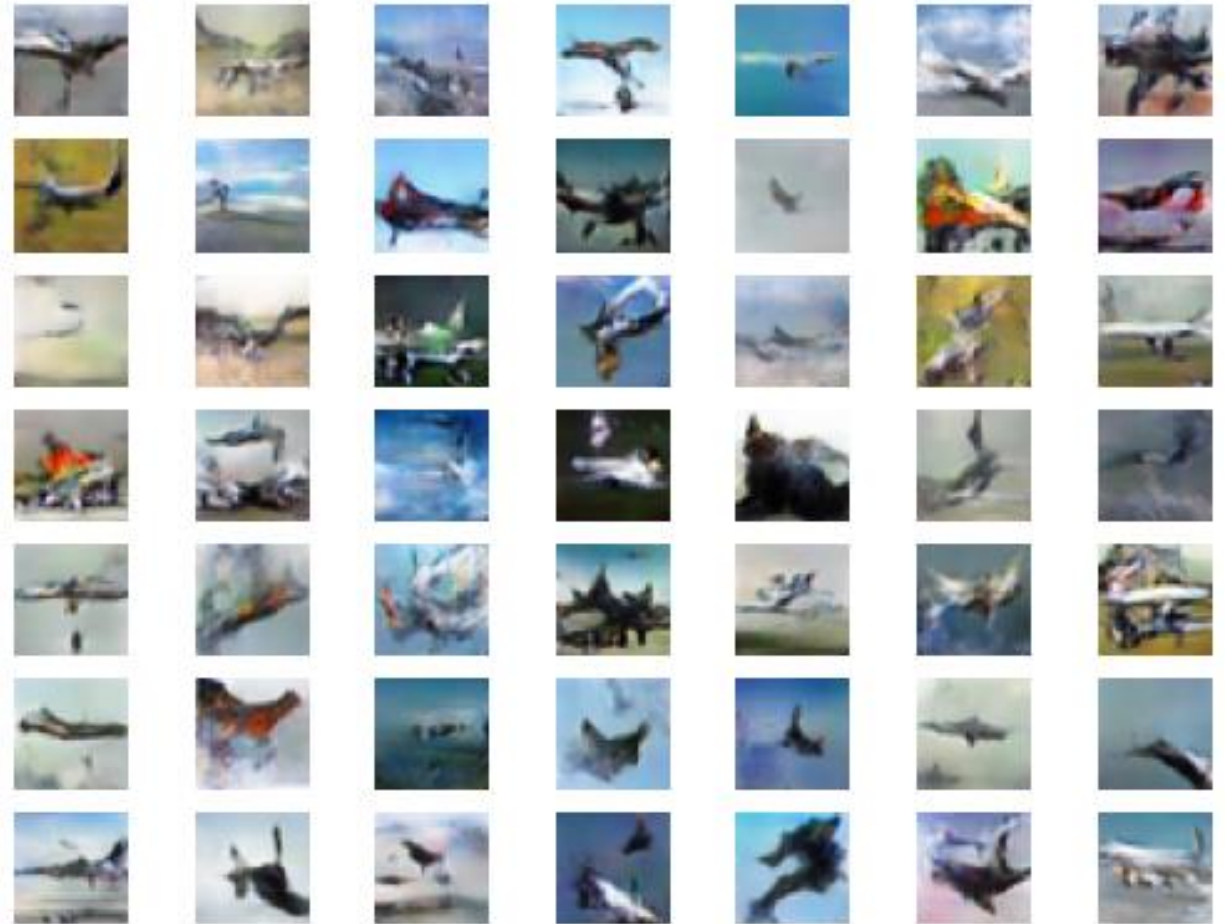
# Evolución del entrenamiento

## □ Época 5200



# GAN - Ejemplo

- En **GAN\_CIFAR10.ipynb** se utiliza una red GAN para generar imágenes de aviones a partir de la clase 0 de CIFAR10.
- Esta es la respuesta del generador después de entrenar la GAN durante 100 épocas.




# Ejemplo: Generando dormitorios



Unsupervised Representation Learning with Deep Convolutional  
Generative Adversarial Networks (2016)

# Redes generativas progresivas

- La calidad generativa del modelo mejora incrementando el tamaño de las imágenes en forma progresiva





# thispersondoesnotexist.com

- [thispersondoesnotexist.com](https://thispersondoesnotexist.com) es una web creada por nvidia que crea caras en alta definición de personas que no existen.
- Ejemplo: 3 caras generadas

